




Histogram of Oriented Gradients and Object Detection

by Adrian Rosebrock on November 10, 2014 in Machine Learning, Tutorials



If you've been paying attention to my Twitter account lately, you've probably noticed [one](#) or [two](#) teasers of what I've been working on — a Python framework/package to rapidly construct object detectors using Histogram of Oriented Gradients and Linear Support Vector Machines.

Honestly, I really can't stand using the Haar cascade classifiers provided by OpenCV (i.e. the Viola-Jones detectors) — and hence why I'm working on my own suite of classifiers. While cascade methods are extremely fast, they leave much to be desired. If you've ever used OpenCV to detect faces you'll know exactly what I'm talking about.



In order to detect faces/humans/objects/whatever in OpenCV (and remove the false positives), you'll spend a lot of time tuning the `cv2.detectMultiScale` parameters. And again, there is no guarantee that the exact same parameters will work from image-to-image. This makes batch-processing large datasets for face detection a tedious task since you'll be very concerned with either (1) falsely detecting faces or (2) missing faces entirely, simply due to poor parameter choices on a per image basis.

There is also the problem that the Viola-Jones detectors **are nearing 15 years old**. If this detector were a nice bottle of Cabernet Sauvignon I might be pretty stoked right now. But the field has advanced substantially since then. Back in 2001 the Viola-Jones detectors were state-of-the-art and they were certainly a huge motivating force behind the incredible new advances we have in object detection today.

Now, the Viola-Jones detector isn't our only choice for object detection. We have object detection using keypoints, local invariant descriptors, and bag-of-visual-words models. We have Histogram of Oriented Gradients. We have deformable parts models. Exemplar models. And we are now utilizing Deep Learning with pyramids to recognize objects at different scales!

All that said, even though the Histogram of Oriented Gradients descriptor for object recognition is nearly a decade old, it is still heavily used today — and with fantastic results. The Histogram of Oriented Gradients method suggested by Dalal and Triggs in their seminal 2005 paper, *Histogram of Oriented Gradients for Human Detection* demonstrated that the Histogram of Oriented Gradients (HOG) image descriptor and a Linear Support Vector Machine (SVM) could be used to train highly accurate object classifiers — or in their particular study, human detectors.

Histogram of Oriented Gradients and Object Detection

I'm not going to review the entire detailed process of training an object detector using Histogram of Oriented Gradients (yet), simply because each step can be fairly detailed. But I wanted to take a minute and detail the general algorithm for training an object detector using Histogram of Oriented Gradients. It goes a little something like this:

Step 1:

Sample P positive samples from your training data of the object(s) you want to detect and extract HOG descriptors from these samples.

Step 2:

Sample N negative samples from a *negative training set* that **does not contain** any of the objects you want to detect and extract HOG descriptors from these samples as well. In practice $N \gg P$.

Step 3:

Train a Linear Support Vector Machine on your positive and negative samples.

Step 4:




Figure 2: Example of the sliding a window approach, where we slide a window from left-to-right and top-to-bottom. *Note: Only a single scale is shown. In practice this window would be applied to multiple scales of the image.*

Apply hard-negative mining. For each image and each possible scale of each image in your negative training set, apply the sliding window technique and slide your window across the image. At each window

compute your HOG descriptors and apply your classifier. If your classifier (incorrectly) classifies a given window as an object (and it will, there will absolutely be false-positives), record the feature vector associated with the false-positive patch along with the probability of the classification. **This approach is called *hard-negative mining*.**

Step 5:

Take the false-positive samples found during the hard-negative mining stage, sort them by their confidence (i.e. probability) and re-train your classifier using these hard-negative samples. *(Note: You can iteratively apply steps 4-5, but in practice one stage of hard-negative mining usually [not not always] tends to be enough. The gains in accuracy on subsequent runs of hard-negative mining tend to be minimal.)*

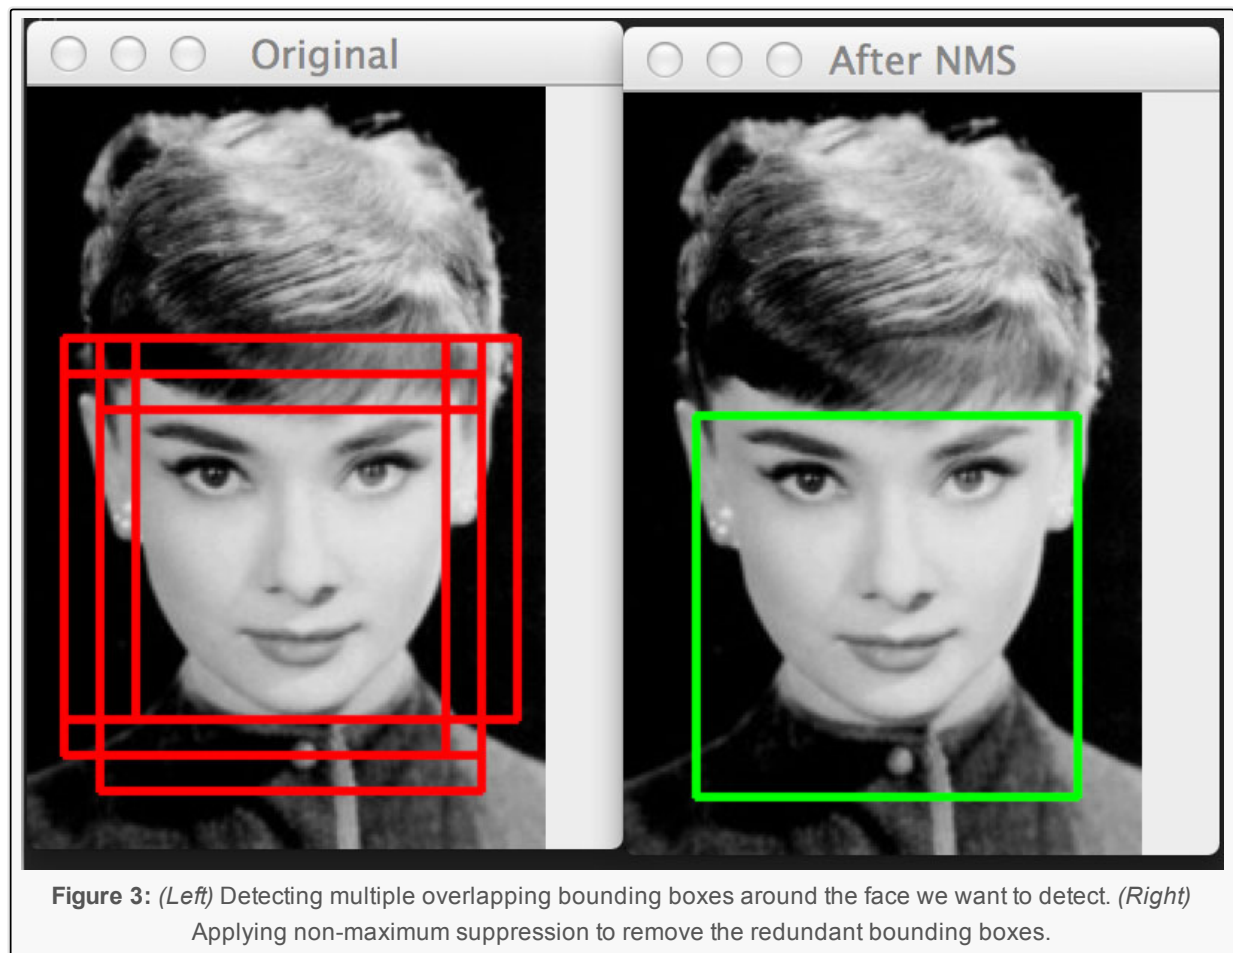
Step 6:

Your classifier is now trained and can be applied to your test dataset. Again, just like in Step 4, for each image in your test set, and for each scale of the image, apply the sliding window technique. At each window extract HOG descriptors and apply your classifier. If your classifier detects an object with sufficiently large probability, record the bounding box of the window. After you have finished scanning the image, apply non-maximum suppression to remove redundant and overlapping bounding boxes.

These are the bare minimum steps required, but by using this 6-step process you can train and build object detection classifiers of your own! Extensions to this approach include a [deformable parts model](#) and [Exemplar SVMs](#), where you train a classifier for *each positive instance* rather than a *collection of them*.

However, if you've ever worked with object detection in images you've likely ran into the problem of *detecting multiple bounding boxes around the object you want to detect in the image*.

Here's an example of this overlapping bounding box problem:



Notice on the *left* we have 6 overlapping bounding boxes that have correctly detected Audrey Hepburn's face. However, these 6 bounding boxes all refer to the same face — we need a method to suppress the 5 smallest bounding boxes in the region, keeping only the largest one, as seen on the *right*.

This is a common problem, no matter if you are using the Viola-Jones based method or following the Dalal-Triggs paper.

There are multiple ways to remedy this problem. Triggs et al. suggests to use the [Mean-Shift algorithm](#) to detect multiple modes in the bounding box space by utilizing the (x, y) coordinates of the bounding box as well as the logarithm of the current scale of the image.

I've personally tried this method and wasn't satisfied with the results. Instead, you're much better off relying on a *strong classifier* with *higher accuracy* (meaning there are very few false positives) and then applying non-maximum suppression to the bounding boxes.

I spent some time looking for a good non-maximum suppression (sometimes called non-maxima suppression) implementation in Python. When I couldn't find one, I chatted with my friend [Dr. Tomasz Malisiewicz](#), who has spent his entire career working with object detector algorithms and the HOG descriptor. There is literally *no one* that I know who has more experience in this area than Tomasz. And if you've ever read any of his papers, you'll know why. His work is fantastic.

Anyway, after chatting with him, he pointed me to two MATLAB implementations. [The first](#) is based on the work by [Felzenszwalb et al.](#) and their deformable parts model.

The [second method](#) is implemented by Tomasz himself for his [Exemplar SVM](#) project which he used for his dissertation and his ICCV 2011 paper, [Ensemble of Exemplar-SVMs for Object Detection and Beyond](#). It's important to note that Tomasz's method *is over 100x faster* than the Felzenszwalb et al. method. And when you're executing your non-maximum suppression function millions of times, that 100x speedup really matters.

I've implemented both the Felzenszwalb et al. and Tomasz et al. methods, porting them from MATLAB to Python. Next week we'll start with the Felzenszwalb method, then the following week I'll cover Tomasz's method. While Tomasz's method is substantially faster, I think it's important to see both implementations so we can understand exactly *why* his method obtains such drastic speedups.

Be sure to stick around and check out these posts! These are absolutely critical steps to building object detectors of your own!

Summary

In this blog post we had a little bit of a history lesson regarding object detectors. We also had a sneak peek into a Python framework that I am working on for object detection in images.

From there we had a quick review of how the Histogram of Oriented Gradients method is used in conjunction with a Linear SVM to train a robust object detector.

However, no matter what method of object detection you use, you will likely end up with multiple bounding boxes surrounding the object you want to detect. In order to remove these redundant boxes you'll need to apply Non-Maximum Suppression.

Over the next two weeks I'll show you two implementations of Non-Maximum Suppression that you can use in your own object detection projects.

Be sure to enter your email address in the form below to receive an announcement when these posts go live! Non-Maximum Suppression is absolutely *critical* to obtaining an accurate and robust object detection system using HOG, so you definitely don't want to miss these posts!

Resource Guide (it's totally free).

Enter your email address below to get my **free 11-page Image Search Engine Resource Guide PDF**. Uncover **exclusive techniques** that I don't publish on this blog and start building image search engines of your own!

DOWNLOAD THE GUIDE!



◆ histogram of oriented gradients, hog, linear svm, nms, non maxima suppression, non maximum suppression, svm

< How to Display a Matplotlib RGB Image

Non-Maximum Suppression for Object Detection in Python >

152 Responses to *Histogram of Oriented Gradients and Object Detection*



SHI Xudong November 10, 2014 at 11:20 pm #

REPLY ↩

Hi Adrian,

Thank you for your great post.

I am currently working on object detection too. My goal is to make the detection faster, because object detectors are usually slow.

In your post, the object detection framework you are using is a 1 model, N image scales. However, there are alternatives, such as N models, 1 image scale; 1 model, N/K ($K > 1$) image scales (FPDW approach); N/K models ($K > 1$), 1 image scale.

I am thinking about creating a unified framework, which can include all these frameworks, but have no idea about the implementation now. Could you please make some comments:

— Is it worth to try?

— how will you do if you are doing it?

References:

1. Benenson, Rodrigo, et al. "Pedestrian detection at 100 frames per second." Computer Vision and Pattern Recognition (CVPR), 2012 IEEE Conference on. IEEE, 2012.

2. Dollár, Piotr, Serge Belongie, and Pietro Perona. "The Fastest Pedestrian Detector in the West." BMVC. Vol. 2. No. 3. 2010.



Adrian Rosebrock November 11, 2014 at 7:07 am #

REPLY ↩

Hi Shi! You are absolutely right, object detectors can be painfully slow, especially when you using the sliding window technique. And when expensive features such as HOG need to be computed, it can really kill performance.

And you are correct, I am utilizing the N image scales model for this framework. However, I have plans to create something similar to the FPDW approach by Dollar et al.

Right now performance using the N image scales model is actually pretty good for my framework, but not great. The reason is because I have distributed the image pyramid to all available cores of the system — this is an obvious solution where making the HOG sliding window computation run in parallel can dramatically speedup the code.

However, doing something like FPDW will further increase the speed (but lessen accuracy slightly).

Send me an email and I would be happy to chat more about implementation details. Thanks again for the comment, there is a ton of great information in here.



Eyy November 10, 2014 at 11:28 pm #

REPLY ↩

Really good illustration. Thanks! waiting for your following posts



Varun Kumar November 11, 2014 at 5:15 am #

REPLY ↩

Thanks Adrian for this and other wonderful posts. I have been following your posts for quite a long time now. I must say you are doing a wonderful job. Keep it up!



Adrian Rosebrock November 11, 2014 at 7:08 am #

REPLY ↩

Thanks, I really appreciate it! 😊



P Derep November 11, 2014 at 8:36 am #

REPLY ↩

La Fin Du Monde... that's quite a good beer! I'm impressed! 😊 Kudos for the interesting article, BTW.



Wajih Ullah Baig November 11, 2014 at 12:24 pm #

REPLY ↩

Good post, waiting for some nifty code to mess with 😊
Your blog is now my standard technology for teaching!



Adrian Rosebrock November 11, 2014 at 2:35 pm #

REPLY ↩

Thank you for such an awesome compliment! 😊



Wajih Ullah Baig November 13, 2014 at 2:53 am #

REPLY ↩

Thank you, not me for such an awesome blog!!



Niv November 17, 2014 at 9:29 am #

REPLY ↩

Nice post, thanks.

In phase #5, the false positives are taken along with their probabilities and then sorted by their probabilities in order to further retrain the classifier. I have two questions about which I would appreciate to get a clarification:

1. why do you need their probabilities in that retraining phase? linear SVM doesn't require their probabilities but rather merely their taggings (negatives in this case).
2. Why do you need them to be sorted? linear SVM doesn't require any sorting of the training samples.



Adrian Rosebrock November 17, 2014 at 10:49 am #

REPLY ↩

You would want to keep track of the probabilities for two reasons. The first being that you may not have enough physical memory to store the positive samples, negative samples, and hard-negative samples and train your SVM. In the case you are limited by RAM, you would want to sort the samples by their probability and keep only the samples the classifier performed "worst" on, discarding the rest. Secondly, you may be worried about overfitting.



asaaki December 18, 2014 at 3:26 am #

REPLY ↩

Will Steps 4 and 5 work for a multiclass cascade classifier as well, that uses a boosting algorithm? I notice in your example you're dealing with LSVMS.



Adrian Rosebrock December 18, 2014 at 7:45 am #

REPLY ↩

Potentially, but there are other approaches that are more robust (and faster to train) than cascade classifiers.



Karthik January 6, 2015 at 1:33 am #

REPLY ↩

According to the paper published by Dalal and Triggs, they suggest gamma correction is not necessary. I have the doubt about whether correcting gamma is a good option to go for or not. If Gamma correction is necessary what is the gamma value I have to take for better performance. Awaiting for reply.!!! thanks in advance!!



Adrian Rosebrock January 6, 2015 at 7:38 am #

REPLY ↩

That is correct, gamma correction is not necessary. Normalization however is quite often helpful. You can normalize by either taking the log or the square-root of the image channel before applying the HOG descriptor (normally the square-root is used). Another method to make HOG more robust is to compute the gradient magnitude over all channels of the RGB or L*a*b* image and take the maximum gradient response across all channels before forming the histogram.



Karthi January 6, 2015 at 2:40 pm #

REPLY ↩

Thank you for your reply!! But why can't we consider converting to grayscale and processing would be the better option?



Adrian Rosebrock January 6, 2015 at 2:52 pm #

REPLY ↩

You can certainly convert to grayscale and compute HOG as well. It's just been shown that taking the maximum response over the individual RGB or L*a*b* channels can increase accuracy in some cases. Your mileage may vary depending on your dataset.



Karthik January 6, 2015 at 3:45 pm #

And we are actually stretching the contrast while we are gamma correcting, so which gamma value you think will be providing the higher performance? gamma of 2?



Adrian Rosebrock January 6, 2015 at 4:02 pm #

Please see my previous comment. I would not recommend using gamma correction. Just a simple log or square-root normalization should suffice.



Sarah April 10, 2015 at 11:34 pm #

REPLY ↩

In opencv the gamma correction can only be set or not set. I wonder what is the valued applied by default. For images which can be taken at night. I believe gamma correction might help?



Adrian Rosebrock April 11, 2015 at 8:55 am #

REPLY ↩

Hi Sarah, thanks for the comment. However, log-based gamma correction normally hurts HOG + Linear SVM performance. The square-root (in most cases) or simple variance normalization is better option.



Karthik January 6, 2015 at 10:22 pm #

REPLY ↩

Ya, I got you. And I would like to thank for your valuable blog !!! really helpfull!!



Morné January 26, 2015 at 2:43 am #

REPLY ↩

Hi Adrian, this is great stuff. Do you have any tutorials on implementing HOG descriptors. Having that is really helpful for learning the essentials. Thanks.



Adrian Rosebrock January 26, 2015 at 6:43 am #

REPLY ↩

Hi Morne, check out the scikit-image implementation of [HOG](#). That is a great starting point!



James January 28, 2015 at 3:08 am #

REPLY ↩

Hi Adrian,

Great post! Are you still working on this or is it already completed?



Adrian Rosebrock January 28, 2015 at 6:28 am #

REPLY ↩

Great question! The framework is 100% complete at this point. I'll be covering it inside the computer vision course I am working on. Check out the [Kickstarter](#) for more information!



Jeff February 12, 2015 at 11:18 am #

REPLY ↩

I have a suggestion for speeding up the search, and would like your opinion.

What if you applied the KPM, or Boyer-Moore to the search ? Meaning, convert the histogram 'coefficients' to a string representation, and do the same with the image being searched. Then, look at the 'suffixes' first.

I've used the Boyer-Moore, because it's about 10x faster for string searching. In this application, it would be like looking at the right edge of a rect first, and then skipping over pixels when there isn't a match.



Jeck March 14, 2015 at 1:39 pm #

REPLY ↩

Hi Adrian, awesome tutorial btw. Im quite a beginner in openCV. I want to create a computer vision algorithm that is able to detect license plates and read them. However, I dont really know where to begin. What tutorial do you suggest that I can start with. So far, I was able to create the document scanner in your tutorial in c++.

Thanks in advance.



Adrian Rosebrock March 14, 2015 at 4:17 pm #

REPLY ↩

Hi Jeck, great question. I'll be covering license plate detection and recognition inside the [PyImageSearch Gurus course](#). License plate detection and recognition is not an easy task, but inside the course I'll be breaking it down into simple and easy to digest components, similar to other tutorials on the PyImageSearch blog.



Miriam March 17, 2015 at 6:10 am #

REPLY ↩

Hi , i'am working at detection face and i have a problem. the algorithm didn't detect the faces in in different position. what's the problem ?

Thank you



Adrian Rosebrock March 17, 2015 at 6:43 am #

REPLY ↩

Hi Miriam. Unfortunately, there could be many, many reasons why faces are not detected and your question is a bit ambiguous. Here are just some example questions for you to consider: Are you using HOG + Linear SVM? Have you chosen the optimal HOG parameters for your descriptor? What about your Linear SVM? Are you using a soft-classifier? Do your training examples sufficiently represent the faces that you want to detect in images? Are you performing hard negative mining to increase accuracy? Are you using image pyramids so you can detect faces at multi-scales? Or using a sliding window so detect faces in different positions of the image? Be sure to consider all avenues.



Miriam April 7, 2015 at 7:26 am #

REPLY ↩

Ok

Thank you Adrian.



Milos March 24, 2015 at 4:33 am #

REPLY ↩

Hey Adrian, i really want to thank you for outstanding lessons you taught me about computer vision.

I was wandering, can we expect something about image stitching. Something like creating panorama images? 😊

All best to you sir! 😊

**Free 21-day crash course on
computer vision & image search
engines**



Adrian Rosebrock March 24, 2015 at 6:36 am #

REPLY ↩

Hi Milos, thank you for such a kind comment! And we will absolutely be doing image stitching and panoramas within the next few months. I can't set an actual date yet since I have a series of posts coming out soon, but I'll be sure to let you know when the image stitching post goes live.



fansari March 25, 2015 at 5:33 am #

if my dataset is on colored images and color should I use grayscale image or should I find gradient the result

**Free 21-day crash course
on computer vision &
image search engines**

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



Adrian Rosebrock March 25, 2015 at 6:48 a

Dalal and Triggs actually recommend computing the color space and computing the gradient magnitude respectively, and then taking the maximum magnitude to improve performance.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Yiwen Wan April 10, 2015 at 11:29 pm #

REPLY ↩

Thanks for your great work. I have a question about training positive and negative samples. They can be of different size. Do you try to resize them to a predefined size? I concern resizing may change the original gradients orientation depending on how it is resized. Looking forwards to your answer.



Adrian Rosebrock April 11, 2015 at 8:56 am #

REPLY ↩

Indeed, each object of each class should be resized to a pre-defined size to ensure the resulting HOG feature vectors are the same dimensionality.

Free 21-day crash course on computer vision & image search engines



Yiwen Wan April 14, 2015 at 7:50 pm #

Thanks. I am aware of that. Just want to make sure whether it will give distortion or not. If it does bring distortion to the gradient orientations. Maybe compensation is needed? Or maybe the distortion is too little to consider.



Adrian Rosebrock April 15, 2015 at 9:40 am #

It will absolutely bring in some distortions since we are ignoring the aspect ratio of the image during the resizing. However, if we can get features of the same dimensionality, we can get features of the same dimensionality.

Free 21-day crash course on computer vision & image search engines



Méabh Loughman April 17, 2015 at 5:01 am #

Hey Adrian ,

Just wanted to say thank you. I am a student studying computer vision and I've been taught masters and this blog really helped me so much.

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

intuitively 😊

Thanks again,
Méabh

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock April 17, 2015 at 7:03 am #

Thanks so much Méabh 😊 I'm happy to help. Best of luck with your masters program!



Rish April 18, 2015 at 9:40 am #

REPLY ↩

First I am just thank you for your wonderful and super easy to understand tutorials and perhaps the best available.

I have one question on Object Recognition using sliding window and SVM. I am using C++ and the SVM on OpenCV. On detection I get multiple windows where I need to apply Non-Maximum Suppress (which I learnt well from your tutorial). However, the linear SVM output is a hard decision of +1 for objects and -1 for non-objects. In this case its not possible to do NMS as all weights (considering the prediction response) are same.

I read a paper from Platt, 1999 to convert the prediction response to probability. However, I am wondering if you know there is any simpler or better way to achieve this perhaps

Thank you for your time.

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock April 18, 2015 at 1:33 pm #

REPLY ↩

Hey Rish, thanks for the comment. Does your SVM library not support returning probabilities? Most SVM libraries do. I don't do much work in C++ (literally none), but I know the scikit-learn SVMs support returning probabilities along with class labels.

Also, you can still do NMS without probabilities. Instead of using the probabilities, just use the bottom right-hand corner, like I do in this post on [fast non-maximum suppression](#). It's obviously better if you have the probabilities, but this approach

Free 21-day crash course on computer vision & image search engines



Rish April 19, 2015 at 11:23 pm #

Hi Adrian, thank you for your prompt response. I was able to get probabilities while SVM in OpenCV C++ does

Regarding the NMS, as your tutorial is implemented, I was confused how to use Bounding Box probabilities

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

wrong than could please guide me on the right

→ Instead of computing the overlap should

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock April 20, 2015

You'll need to modify my NMS code to accept a list of probabilities that corresponds to the bounding boxes. Then you'll need to sort those probabilities (Line 27) rather than the bottom-right corner.



Rish April 23, 2015 at 12:42 pm #

Hi Adrian. Thank you very much I managed to get it working with LibSVM. However, I noticed that getting the probabilities from SVM is computationally more expensive compared to the label prediction (hard decision). I wonder if this happens or did I just got something wrong? Mathematically, I think estimating probability needs more operations. However, my question here is that if I don't use probability but just use label predication i.e (+1/-1) than use probability, do you think that is also good enough or Probabilities are not needed for SVM results? (Apologies my poor English).

Secondly, about the HOG descriptor. Say for example I trained a classifier to detect standing bottle (say sliding window of 64×128) = 3780 dimension feature vector. During detection, I will use this window size to detect the bottle. However, my question is when the bottle is resting on its side (now say the dimension is 128×64) but our detector uses the window of 64×128 , how do I cater for this issue? I read paper on Part based model but I think PBM is computationally expensive. Perhaps I thought to divide the training image into 4 parts (say 16×32) and train this. Do you have any suggestion on this?



Adrian Rosebrock April 23, 2015 at 1:08 pm #

Hey Rish, you're absolutely right about the computational cost. However, if you use a more precise final bounding boxes. It really depends on the tolerance you're willing to accept.

As for your second question, HOG based on the fact that there is a certain "structure" to the image. For example, to recognize a bike, it would look for two wheels. If the wheels were rotated 90 degrees, you would need to use different descriptors and hence in your case you would need to rotate the object is rotated.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to

If you're really concerned about finding the right input on their sides, just take each input in turn. In this case, you'll be evaluating each image four times, increasing the computational cost, but you'll find the right

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Rish May 4, 2015 at 12:54 pm #

Hi Adrian. Sorry I was busy with my Finals so was unable to complete my work. Actually, I have got most of the HOG detection implemented in C++. As I mentioned above, that I want to change your NMS from Bottom Right Corner to probability sorting as you mentioned.

I think I am bit lost, I removed all the bounding box params from your NMS and use the Probability instead I am confused like in line 51 where you done the overlapping computation. How can I change this to probability instead? Do I need to do something is Gradient Decent? Sorry I am just confused, please help me out. Thanks



Adrian Rosebrock May 4, 2015 at 1:15 pm #

No need for gradient descent or anything fancy. When you make a call to your NMS function you'll need to pass in three variables: the bounding boxes, the overlap threshold, and a probability of detection that corresponds to the bounding box. Then you need to sort (from highest probability to lowest probability) the bounding boxes based on their probability. From there, everything will be the same. Like I said, the only code that you'll need to change is the code that sorts on probability rather than the bottom-right coordinate.

Free 21-day crash course on
computer vision & image search
engines



Siva Krishna January 30, 2016 at 5:08 am #

Hi Andrian,

I have few doubts regarding NMS algorithm using probabilities

1. What does probability means? It is not clear to me.
2. You said that we have to sort in decreasing order, since we are interested in the highest probability, we have to sort in increasing order, since we are interested in the lowest probability.

Thank you for your wonderful posts.

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



Adrian Rosebrock January 30, 2016 at 5:08 am #

The probabilities actually come from the sorting of the probabilities, that depends on the NMS algorithm.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Russell May 2, 2015 at 3:24 pm #

I am new to sklearn. I see many of you have advanced experience in this area. I am therefore asking for your assistance so that I can catch up 😊

I have a number of images of sedans, SUV's and Trucks that pass outside my house. I want to train a SVM classifier to detect whether a vehicle is a one of these three. I have cropped all of the images and successfully created Histogram Oriented Gradient images of these vehicles. My question is Do I have to write code to reshape the image and format it to get the data into a format useable by scikit-learn or is there code already written to do this. I am grateful for your help.



Adrian Rosebrock May 2, 2015 at 3:36 pm #

REPLY ↩

Visualizing a Histogram of Oriented Gradients feature vector. While you can visualize your HOG image, this is not appropriate for training a classifier — it simply allows you to visually inspect the gradient orientation/magnitude for each cell. Instead, you'll need to extract the HOG descriptor using something like [scikit-image](#) and then pass the feature vector into scikit-learn. As I mentioned in an email to you, I'll be covering all this inside the [PyImageSearch Gurus course](#).

Free 21-day crash course on computer vision & image search engines



Russell May 2, 2015 at 3:59 pm #

REPLY ↩

Hi Adrian,

Yes, I understand that the HOG image is not useable for integrating into Scikit learn. I understand that it needs to be reshaped into a matrix of row vectors so that we have number of samples vs number of features. Are you telling me that this is done in scikit-image? Thanks for your help.



Adrian Rosebrock May 2, 2015 at 4:01 pm #

Yes, take a look at the [documentation](#). The Handwriting Recognition chapter of Case Study 1: Document Image Classification shows how to extract a feature vector. Secondly, the HOG image is essentially useless for anything but visualization (which scikit-image provides you with) to

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



Jay July 29, 2016 at 7:29 am #

Hi Adrian, it occurred to me I should have looked into scikit-image's source code and vote for histogram orientation binning value falling into which bin, which is not quite an accurate way for computing hog.

Actually, I found in my experiments that object detection based on OpenCV's implementation is more accurate than with the one from scikit-image, and runs faster.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock July 29, 2016 at 8:24 am #

Thanks for sharing your experience Jay! I've been using a personal modification of the scikit-image HOG that I've been meaning to create a pull request for. It handles various types of input transformations (square-root, log, variance) along with multi-channel images, allowing you to take the maximum of the gradient across the channels (this tends to work better in practice). I'll have to look into weighted vote binning as well.

Free 21-day crash course on
computer vision & image search
engines



TimothyUntan May 5, 2015 at 9:47 pm #

Dang, you're awesome man... That one helps a lot for my final year project... Cheers...



Adrian Rosebrock May 6, 2015 at 7:06 am #

REPLY ↩

I'm happy it helped Timothy! And congrats on your (impending) graduation!



pal May 7, 2015 at 5:30 am #

great post!

I just want to know that how many training images are



Adrian Rosebrock May 7, 2015 at 7:16 am #

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

The answer is, it depends! Some problems are simple, some are very complex. Other problems are very, very complex. To start with at least 100 positive training examples and 1000 negative training examples. And then perform 1000 false positive examples. This approach will give you a baseline to further tune your parameters. don't expect your first attempt to give you the best results.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Pal May 8, 2015 at 12:06 am #

REPLY ↩

Thanks a lot! 😊



abbas May 22, 2015 at 12:48 am #

REPLY ↩

Hi Adrian,

I am a student of BS computer science i have started working on human detection in image using HOG descriptor. I have a problem in testing phase step 4 and 5. Are you telling me what i do i use a matlab as a tool.

Free 21-day crash course on
computer vision & image search
engines



Adrian Rosebrock May 22, 2015 at 5:09 am #

REPLY ↩

Hi Abbas, I honestly have not used MATLAB in a very, very long time so I am probably not the right person for that question.



ALI June 25, 2015 at 4:59 am #

REPLY ↩

Thank you for your work , May I ask you to direct me to some implementation of hard-negative mining. Actually I didn't get the point of how to reuse the false negative data ! . Should I use multi instance Svm ; I mean like clustering the false positive .

In case I get false positive by my trained classifier on train data ? or keep them with negative samples ? I should be taken during retrain the classifier , May you

Cheers ,

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



Adrian Rosebrock June 25, 2015 at 6:19 am #

So basically, hard negative mining is all about finding HOG feature vectors) that the Linear SVM gets to pretend we are training an object detector to detect. We extract HOG features from a positive dataset (the face dataset) and a negative dataset (which contains absolutely no faces), then train our Linear SVM.

However, this classifier may falsely detect faces in images where there are no faces. So, we run our classifier on the negative data (which contain no faces what-so-ever), and we collect any HOG feature vectors that the classifier incorrectly reports as a face. This process is called “hard-negative mining”. In your case, if you get a false-positive from your trained classifier you want to **keep** that data since you can use it to better improve your classifier.

Finally, we need to re-train our classifier, which is just a 1-vs-1 SVM: either “face” or “not a face” using the HOG feature vectors from the original face dataset, the HOG feature vectors from the non-face dataset, as well as the hard-negative HOG feature vectors.

We then let this classifier train, with the end goal (normally) being that our classifier better detects faces, while ignoring non-faces.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



vin July 3, 2015 at 2:15 pm #

hi adrian!

in your view, do you think the success of convolutional NNs have made other image processing techniques (eg hog + svm) obsolete? thanks!

Free 21-day crash course on computer vision & image search engines

REPLY



Adrian Rosebrock July 3, 2015 at 4:10 pm #

Definitely not. CNNs, and deep learning in general, is a tool just like anything else. There is a time in a place for it to be used. In some situations it's warranted. In some situations it's overkill. And in other it's just the flat out wrong choice. Take a look at my post on the [deep learning “band wagon”](#) mentality that rises up every 5-7 years.

REPLY



ngapweitham July 7, 2015 at 8:02 am #

Hi, I have two questions after reading this post

1 : Would you show us step 1~step 6 in your blogs just step 6(<http://www.pyimagesearch.com/2014/11/17/non-max-suppression-for-object-detection/>) or no step 1 ~step 5.Or I have to wait until the course “

2 : The SVM of opencv do not provide probability estimates, should I replace SVM?Or I should find another SVM libraries

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to

Again, thank you for your brilliant posts



Adrian Rosebrock July 7, 2015 at 8:47 am

1. I plan to review Steps 1-6 inside the [PylImageSearch Gurus](#) steps on the blog later on, but right now I am planning on detailing every step inside PylImageSearch Gurus.

2. I would recommend using the [scikit-learn implementation of the Linear SVM](#).

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



ngapweitham July 7, 2015 at 11:50 am #

REPLY ↩

Thanks for your suggestion, I have some questions to ask(new to the field of machine learning)

1 : what are the pros and cons of using the scikit-learn linear svm vs the random forest of opencv3.0? I think the random forest of opencv provide probability estimation too(<http://stackoverflow.com/questions/28303334/multi-class-image-classification-with-probability-estimation>)?

2 : Since HOG is a feature descriptor(am I right?), is it possible to use another descriptor to describe the object and feed to svm? Like akaze, kaze, brisk, freak(In truth, I do not know their different) and so on

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock July 7, 2015 at 1:31 pm #

REPLY ↩

Simply put, a Linear SVM is very fast. A random forest classifier is capable of generalizing to non-linear spaces; however, it requires computing N number of trees. It also doesn't work well for sparse feature spaces which Linear SVMs excel at. SVMs can also return probability estimations as well.

2. HOG is indeed an image descriptor. You could certainly use any other image descriptor and feed it into your classifier, such as Local Binary Patterns, Zernike moments, Haralick texture. And yes, you can AKAZE, BRISK, SURF, SIFT. I'm covering all of these topics in detail inside the [PylImageSearch Gurus](#). I've also played around with machine learning, [Practical Machine Learning](#)

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to



ngapweitham July 10, 2015 at 6:17 am #

According to the dlib document(http://dlib.net/fhog_0) to detect semi-rigid object, I have two questions

- 1 : what is semi-rigid object?
- 2 : smoke is rigid-object or not?should I use HOG + S

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock July 10, 2015 at 6:27 am #

REPLY ↩

A semi-rigid object is something that has a clear form and structure, but can change slightly — consider how a human walks. We put one foot in front of the other, maybe move our arms a bit. The form is clearly the same, but it does change a little bit. HOG can be used to detect semi-rigid objects like humans provided that our poses are not too deviant from our training data. Finally, smoke is definitely *not a semi-rigid object*. *Smoke diffuses into the air and has no real shape or form.*



ngapweitham July 10, 2015 at 11:02 pm #

REPLY ↩

Thanks for your reply.What kind of descriptors you will recommend if you want to classify smoke?

The thesis “Forest smoke detection using CCD camera and spatial-temporal variation of smoke visual patterns” use random forest

Another one “SMOKE DETECTION USING TEMPORAL HOGHOF DESCRIPTORS AND ENERGY COLOUR STATISTICS FROM VIDEO” use HOG + HOF(I do not what is this yet)

I haven't go through the details of these papers yet, the second paper claim that the HOG can separate rigid object and non-rigid object very well.They use HOF to estimate the motion of the smoke, but it would be a problem if the video is not stable

Some implementation(Video-based Smoke Detection: Possibilities, Techniques, and Challenges) even do not use machine learning at all

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock July 11, 2015 at 6:3

I don't have any experience with sm
I would apply something like a spatial pyramid pooling
the L*a*b* and then I would create an exper
such as Local Binary Patterns (I would be re
texture, and I would even try HOG to see ho
that's why we perform experiments.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



ngapweitham July 11, 2015 at 8:23 am #

Thanks for your suggestion, I will try to find

After studying the cases of Plant classification, Hand learning(biased/overfit), I have more sense about w

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Douglas July 20, 2015 at 9:12 am #

REPLY ↩

Hi Adrian,

Awesome job, your posts are the best posts about OpenCV that I found!! Congratulations!!

I have one question: Did you ever worked with SVM for ONE_CLASS? I tried to find something about it but got nothing.



Adrian Rosebrock July 20, 2015 at 10:14 am #

REPLY ↩

Hey Douglas, thanks for the comment. learning implementations. I find the ones implem as customizable.

Free 21-day crash course on computer vision & image search engines



Marine July 27, 2015 at 8:02 am #

REPLY ↩

Hi Adrian,

Thank you for this great tutorial !

I'm trying to train a svm for a face detection system. I've got positive and negative examples.

Do you know the size I should choose for the hog extraction ? I know it's often 64*128 for human detection, with blocksize =8. But I don't know for face problem.. I'm trying to calculate hog features on 25*35 images, with the function hog.compute(...) but it's not working...

Any help ? 😊

Plus, how do you know the optimal parameters for s

Thanks !

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock July 27, 2015 at 10:17 am #

You normally would experiment with the goes for the parameters of the SVM. I actually pr

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

HOG and the [scikit-learn implementation of the](#) cross-validation to find the optimal values.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Ahmad August 7, 2015 at 6:46 am #

Hi Adrian,

As a research problem, I want to apply HOG for Content based image retrieval and definitely it is computationally expensive for large datasets. Please guide what to implement and test for efficient image retrieval. I've been tested the effects of changing block, cell, bin, orientation angel etc. , it give a bit performance gain but how does is this sufficient for defense?

Thanks,
Ahmad



Adrian Rosebrock August 7, 2015 at 7:04 am #

REPLY ↩

Hey Ahmad — can you let me know what the computational bottleneck is for your particular project? Is it the feature extraction itself? Or is it the actual search that is slow? If it's the search, it's probably because you are doing a linear scan of the dataset (i.e. comparing the query features to every other feature vector inside the dataset). If at all possible, I would suggest using an approximate nearest neighbor data structure (such as [Annoy](#)) to speed up the actual search.

Free 21-day crash course on computer vision & image search engines



Ahmad August 7, 2015 at 9:53 am #

REPLY ↩

Thanks Adrian for very quick reply.

Actually I'm pursuing my PhD. For my research contribution the above query was asked. In my last effort I used HOG with PCA and classify using SVM at Corel dataset. Then I tuned HOG with its parameters like bin,angel,block size etc. but it is not a contribution. I'm stuck at how can I add my contribution.

Ideas:

- > To focus 'rotation invariance' in HOG, use
 - > What to apply to achieve more efficiency ,
- New your valuable input on it!

Thanks,
Ahmad

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



Adrian Rosebrock August 8, 2015

Nice, congrats on working on y
HOG is definitely not rotation invariant. In
some level of rotation invariance. I don't
combined both HOG and LBP together, a
the topic — it might be worth doing some
an experiment to see what your results look

Enter your email below to start your journey to
becoming a computer vision master.

LET'S DO IT!



vin September 14, 2015 at 11:08 am #

REPLY ↩

hi adrian,

what do you think is the next evolutionary step for deformable parts model? it seemed like it was the
darling of the cv community 3-4 years ago...



Adrian Rosebrock September 15, 2015 at 6:09 am #

REPLY ↩

The next evolution is already here and (no pun intended), it's called Convolutional Neural
Networks. CNNs are very accurate for image classification. However, the
largest drawback is the incredible amount of data it takes to train supervised machine learning algorithms require
supervised machine learning algorithms require a lot of data, but CNNs are particularly data
hungry. In those cases where data is minimal and data augmentation is not a possibility, HOG-
based methods still work very well.

**Free 21-day crash course on
computer vision & image search
engines**



Chi November 4, 2015 at 5:42 am #

REPLY ↩

Thank you for this great tutorial !



Adrian Rosebrock November 4, 2015 at 6:29 am #

REPLY ↩

Thanks Chi! 😊



Philippe Dykmans December 1, 2015 at 5:00 pm

Hi Adrian,

While I totally agree with you that the classifiers provided
immediately follow in the conclusion that this has to do with
simply says that those cascades were (very) poorly trained
with mere patches of a face, rather than a properly trained

**Free 21-day crash course
on computer vision &
image search engines**

Interested in computer vision and image search
engines, but don't know where to start? Let me
help. I've created a free, 21-day crash course that
is hand-tailored to give you the best possible
introduction to computer vision. Sound good?
Enter your email below to start your journey to

images had poor alignment to start with. They probably used the training program that these were faces... Yeah, but OpenCv merely serve as demonstration material, and not as proof. A properly trained Viola-Jones detector, however, can detect faces in images.

That said; I'm interested to test your Python framework.

Gratz,
Philippe

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Rahul December 30, 2015 at 12:26 am #

REPLY ↩

Hi Adrian,

Will it produce good (accurate) results, If we use HOG features and SVM for vehicle licence plate detection?

I only need to detect the rectangle region where license plate is present, recognition is not necessary.



Adrian Rosebrock December 30, 2015 at 6:59 am #

REPLY ↩

It really depends on your license plate dataset. If the license plates are skewed or rotated, it will be harder for HOG + Linear SVM to return good results (since HOG is not rotation invariant).

Free 21-day crash course on computer vision & image search engines



Nicolò January 9, 2016 at 7:43 pm #

REPLY ↩

Hi Adrian,

Thank you for all your great tutorials!

I'm trying to implement Exemplar SVMs with some friends for a school project but we got stuck. We followed your tutorials but our classifier doesn't detect anything. It seems that each single classifier doesn't train properly: as the number of negative sample grows, we get only 2 support vectors (could it be a problem?). We tried `linearSVC(loss='hinge')`, `SVC(kernel='linear')` and `linear_model.SGDClassifier()` but nothing changes. Maybe we're doing something wrong with features extraction? Any hint?



Adrian Rosebrock January 10, 2016 at 7:44 am #

For help with Exemplar SVMs, I suggest you read the paper I authored the original paper on Exemplar SVMs. When you read through the original paper you'll see how to tune the support vectors for each classifier. I have

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

it seems like this could also cause a problem with
extraction is your problem without knowing your

Enter your email below to start your journey to
becoming a computer vision master.

LET'S DO IT!



sruthi February 22, 2016 at 8:46 am #

sir,

i read all the above information regarding HOG.i am currently working on HOG for human detection.do
you have any matlab codes for HOG with which i can go on with the project.
thankyou in advance



Adrian Rosebrock February 22, 2016 at 4:19 pm #

REPLY ↩

No, I do not have any codes related to MATLAB for HOG detection, only Python. You can
find my Python implementation of this HOG + Linear SVM framework inside the [PyImageSearch
Gurus course](#).



Manez February 22, 2016 at 1:45 pm #

Hi Adrian ! Awesome tutorial !

I am using the HOG descriptor from skimage to create a traffic stop sign detector but the process is too
long (approximately 0.5s for a (255, 496) frame).

Did you use the same descriptor as me or the one from Opencv-python ?

Do you have some tips for increasing the speed ?

Free 21-day crash course on
computer vision & image search
engines

REPLY ↩



Adrian Rosebrock February 22, 2016 at 4:17 pm #

REPLY ↩

I use the same scikit-image descriptor as well. As for increasing speed, try splitting
processing of the image pyramid to different cores of the processor, that way each core can
process a separate layer of the pyramid indepen



Manez February 23, 2016 at 9:26 am #

Ok !

I have registered to the PyImageSearch Gurus course.
Are you going to talk about multiprocessing in the future?

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search
engines, but don't know where to start? Let me
help. I've created a free, 21-day crash course that
is hand-tailored to give you the best possible
introduction to computer vision. Sound good?
Enter your email below to start your journey to



Adrian Rosebrock February 23, 2016

I don't directly cover it in the code, but I have seen the source code for implementing a sliding window and share more resources to make multi-pro

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Shahab March 9, 2016 at 6:33 am #

REPLY ↩

Hi Adrian,

Thank you so much for this perfect step-by-step instruction. I'm working on recognizing traffic signs using SVM + HOG. The algorithm is great and the accuracy is very reasonable; however, the speed is deriving me crazy. For each image, of hundreds, it takes more than 5 minutes to yield the results. I really need to shorten this time to at most a minute or less. The sliding window is the main problem. How can I handle the slow moving of sliding window?

Thank you in advance.

Free 21-day crash course on
computer vision & image search
engines

REPLY ↩



Adrian Rosebrock March 9, 2016 at 4:39 pm #

The sliding window itself isn't the actual problem — it's the time associated with extracting the HOG features and then passing them through the SVM. Your pyramid scale may also be hurting performance. I would suggest increasing the size of the image pyramid along with the step size of the sliding window as much as possible without hurting accuracy. The larger the pyramid scale and the sliding window step size are, the less window need to be evaluated.



Dmitry Zaytsev April 19, 2016 at 4:49 am #

REPLY ↩

What about training a Haar (well, LBP) cascade first, lowering false-negatives rate and not bothering much about false-positives. Then, instead of sliding the window just go through every object detected by cascade.

I didn't tried that myself, but maybe you did?

Free 21-day crash course
on computer vision &
image search engines

×



Adrian Rosebrock April 19, 2016

There isn't anything "wrong" with training another classifier and tune any-and better to invest your time in creating a m computation to the GPU. Another appro

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

which is the process of iteratively finding
than applying an exhaustive sliding window

Enter your email below to start your journey to
becoming a computer vision master.

LET'S DO IT!



romanzo April 27, 2016 at 3:11 am #

Hi Adrian great post!

I'm trying to build my own SVM and detect object them on test image using `hog.detectMultiScale`. I tried to build a SVM classifier from sklearn and opencv using crossvalidation. It gives a good accuracy during the cross validation and also while using other test images (however the opencv one just gives me one support vector which is kinda of weird).

If i use sklearn SVM, I load the `coef_` attribute of my SVC object (which corresponds to the primal sv) using `setSVMDetector` but the detection gives a lot of false positive and not even a true positive on the images that were well detected with the sklearn predict...

Have you trained your own svm model and used it with `detectMultiScale`? Any ideas of what I could do wrong.

Thanks a lot!



Adrian Rosebrock April 28, 2016 at 3:23 pm #

I personally don't like using OpenCV to train a custom HOG detector from scratch. Instead I prefer using scikit-learn as it gives me more control. I detail the general process in the post we're currently commenting on and then [demonstrate how to implement HOG + Linear SVM inside PyImageSearch Gurus](#).

Free 21-day crash course on
computer vision & image search
engines

REPLY ↩



Siva Krishna May 14, 2016 at 11:38 am #

Hi Andrian,

I am working on an object detector for fallen people using HOG. As a part of that I have collected some data manually and used some data available online. But I am not obtaining good accuracy. However my data set size is of size train(fall – 400 neg – 1031), test(fall – 246 neg – 503). After initial training I did hard mining also. I have used HOG parameters – 128×64 Bins – 9. How to improve accuracy?. Accuracy doesn't it? Should I play with cell size and other parameters?

Free 21-day crash course
on computer vision &
image search engines



Adrian Rosebrock May 15, 2016 at 11:04 am #

In general, the more data you have the more accurate your model will be. Increase this dramatically. Ideally you should ~1000

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

especially for HOG. Other than that, you will need less of a concern until you can get more data.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Rabelani Netshifhire May 17, 2016 at 4:34 pm #

Hi Adrian. Where is the actual implementation and Code for HOG + SVM, i see in this post it was supposed to be done next week



Adrian Rosebrock May 19, 2016 at 6:11 pm #

REPLY ↩

Sorry for any confusion, but the actual implementation of HOG + Linear SVM is [inside the PyImageSearch Gurus course](#).



Bill May 19, 2016 at 5:39 am #

REPLY ↩

Is it possible to define a distance function between two sets of descriptors for two images? My application is just to iteratively find interesting, similar images, without seeing duplicates.

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock May 19, 2016 at 5:59 pm #

REPLY ↩

Absolutely! Typically you would use either the Euclidean distance or the chi-squared distance. Both are implemented in NumPy/Scipy.



Bill May 23, 2016 at 5:43 am #

REPLY ↩

Given two sets of descriptors, one interesting metric might be the https://en.wikipedia.org/wiki/Hausdorff_distance which would be the worst closest match distance between set features, or maybe the



Adrian Rosebrock May 23, 2016

For comparing histograms, I would recommend using the chi-squared distance — this will give you reasonable performance

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to



Michael June 29, 2016 at 7:41 pm #

Thank you Adrian

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

REPLY ↩



vani July 21, 2016 at 3:10 am #

hi ,

feeling somewhat hard to learn..
trying to implement the code..
a hats off to you..



faryad September 11, 2016 at 2:39 pm #

Thank you Adrian it's very good
Good Luck

REPLY ↩



Romanzo September 19, 2016 at 2:15 am #

Hi Adrian,

In your post you said number of negatives samples >> number of positives samples for the training. I did a bit of research on internet and usually people tend to use balanced data for a binary classification. Could you explain why you are using unbalanced data in this case?

Free 21-day crash course on
computer vision & image search
engines

REPLY ↩



Adrian Rosebrock September 19, 2016 at 1:02 pm #

REPLY ↩

Indeed, that is a great question. Normally we would use a balanced dataset for classification. So why do we use so many more negative examples in the HOG + Linear SVM framework?

The reason is false-positive detections. If we don't have enough negative examples, the classifier will learn to detect false-positive detections.

Free 21-day crash course
on computer vision &
image search engines



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to



Romanzo September 20, 2016 at 3:37 am #

Thanks for the answer. I asked because unbalanced data could actually be the cause of some of

First I'm doing cross validation to confirm i have good data, i get a low FNR and FPR during the training. In the video I get much more FP.

Does that mean my model is overfitting during training?

Secondly, I'm trying to reduce this number of false positives by increasing the number of hard negative samples N and the value of C via random search. But usually the lowest FNR and FPR I get is using 0 hard negative. Which is a bit of a surprise to me.

Does that mean my model was really overfitting and adding some more examples is worst. Should i base my choice of optimal C and N by testing another set rather than getting a FNR and FPR on my training set during cross validation? Any other suggestions? Thanks for your time!

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock September 21, 2016 at 2:14 pm #

REPLY ↩

It's hard to say without seeing your datasets. I would start by asking yourself if your training data is representative of your testing data. If your training data doesn't look anything like your testing data then you can expect to get strange results.

As you suggested, it's normal to set your number of samples and SVM value of C based off a random search or grid search — this is a pretty standard procedure. Given that you are getting the smallest FPR rate when zero hard-negatives is used leads me to believe that your training data might not be representative of your testing data.

Free 21-day crash course on computer vision & image search engines



Romanzo September 27, 2016 at 3:02 am #

For the positive training set, i'm using images from a scene view by a fixed camera. The test images are also taken from this camera. But the negative training set, I'm using the one from INRIA. Should a negative training set be whatever scene without a positive instance or should it reflect the scene (I'm working indoor so should my negative set be only a view from the camera without my object)?



Adrian Rosebrock September 27, 2016 at 3:02 am #

It really depends on where you are in your production environment. It's very common to have an environment that contains objects that you do not want to detect. Your environment will have a lot of varying objects. Using negative views from your camera can help.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

**sonu** September 23, 2016 at 12:39 am #

hi

Which pyramid should be used while implementing w
mean here is gaussian and laplacian etc

Enter your email below to start your journey to
becoming a computer vision master.

LET'S DO IT!

**Adrian Rosebrock** September 23, 2016 at 6:47 am #

REPLY ↩

If review image pyramids for object detection [in this blog post](#). You should actually stay away from blurring your images in the image pyramid when applying the HOG descriptor was that will decrease your accuracy.

**Optimus1072** September 24, 2016 at 6:16 am #

REPLY ↩

Bro! I don't like this site. I come here to read about one topic and ends up spending an hour or two reading different articles 😊

Free 21-day crash course on
computer vision & image search
engines

REPLY ↩

**Gabriel** October 6, 2016 at 2:35 pm #

Hey man,

Great site. I have a question for you. If you train your svm with a certain kind of object size in mint (say 60X180) , then how does the detectMultiScale method figure out what kind of crops it needs to take from the image input so that the descriptor from the compute method will be a vector that's the same size as the ones used in the training process ?

Thanks,
Gabriel

**Adrian Rosebrock** October 7, 2016 at 7:31 am #

REPLY ↩

If you want to train your own custom ob
wouldn't recommend using OpenCV. I would inst
[PyImageSearch Gurus course](#). The Python + Op
also don't have much control over the training pr

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to

**Gabriel** October 7, 2016 at 8:30 am #

I just found out today the hard way that sets svm object. At this point it seems easier to ju and pyramid images and use that with a skel again that i really appreciate the info on this

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock October 11, 2016 at 1:18 pm #

REPLY ↩

No problem Gabriel, I'm happy to help. I know I already mentioned this so I don't want to beat it to death, but I do demonstrate how to train your own HOG detector using sliding windows, image pyramids, and scikit-learns SVM module inside the [PyImageSearch Gurus course](#). If you get stuck implementing on your own (or would like to learn from well documented tutorials) then I would definitely suggest taking a look at the course.



Hoang October 17, 2016 at 11:51 am #

REPLY ↩

Hi Adrian,

Great site for OpenCV and Image Processing.

I have some questions:

1. As you said in Step 2: Sample N negative samples from a negative training set that "does not contain" any of the objects you want to detect and extract HOG descriptors from these samples as well. In practice $N \gg P$.

So example if I want to train a smile detector, the positive images contain many smiling faces and the negative are not smile faces. It's mean that both positive and negative are include face, so it contradict with: "does not contain" ?

2. In case of smile detector, I want to make a classifier with 3-classes: normal – opening mouth (not smile) – smile. I used Viola-Jones algorithm and it's not good. Now, i'm using Logistic Regression (opencv) to make 2 binary classifier: normal – opening mouth, normal – smile. It's better than Viola-Jones, but it still get many false-positives. So in your experience, is that HOG+Linear SVM better? Or could you suggest any approach?

3. With HOG + Linear SVM Model or any Model you suggestion, can I save this model to re-use it in mobile, example save model as xml file and load it in Objective-C (iOS) or Java (Android)? So I'm just predict later, not training again. Finally, may I run re

Free 21-day crash course on
computer vision & image search
engines



Adrian Rosebrock October 17, 2016 at 3:58

1. Semantically this is not a contradiction frowning lips as negative samples. HOG is good lips and look similar. I would experiment with both performance having the detectors trained on data

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

set. Overall, I would suggest you take a look at [this](#). The geometry of these points might help you build a better model.

2. HOG + Linear SVM tends to majorly beat out other models and false-positives. If you're getting many false-negatives, try hard-negative mining.

3. I'm not sure about this since I don't do much mobile development, but you can train a model on a library specific to iOS/Android, but once it's trained, you can deploy it into the production devices.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Walid Ahmed November 4, 2016 at 1:32 pm <#>

REPLY

Thanks a lot



Saloni Mittal November 10, 2016 at 7:12 am <#>

REPLY

Hi Adrian,

I am doing a project on traffic sign detection. Ideally I want to extract HOG features from each image for accurate results if the dataset contains about 2000 positive images and 5000 negatives? Also which svm kernel is preferable?

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock November 10, 2016 at 7:29 am <#>

REPLY

The number of extracted HOG features is entirely dependent on your dataset. A good rule of thumb is to have approximately 5-10x as many negatives as your positives and then perform hard-negative mining. As far as a kernel goes, a linear kernel tends to be standard due to how fast it is. Given your current dataset I would suggest performing an initial experiment to obtain a baseline. From there you can try adjusting the number of negatives and hard-negatives.



Gerardo Rosiles November 10, 2016 at 2:07 pm <#>

Hi Adrian,

Regarding the detection performance. How do you measure the performance of classifiers using metrics like precision and recall?

It seems to me you need to first evaluate the classifier's ability to measure the discriminative ability of the classifier and then try other way.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

My impression is that the detector feeds a very large giving you lots of true negatives and a few (or none) the actual object (which can then be processed with metrics are difficult to apply here, even for a single negative test vectors.

So, I guess my question is, how do you evaluate your best?

Thanks,

Gerardo

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock November 14, 2016 at 12:20 pm #

REPLY ↩

Hey Gerardo — I actually cover the answer to that question in my latest blog post. You use the [Intersection over Union](#) metric to measure the quality of a detection.



cam nguyen November 12, 2016 at 11:02 am #

Free 21-day crash course on
computer vision & image search
engines

REPLY ↩

Hi Adrian, thank you so much for your great article!

I'm learning about human detection using HOG descriptor + Linear SVM.

I am using Verilog HDL to build my system on FPGA board.

My huge issue is : how to train the input dataset by using HOG descriptor? The detailed steps?

Because verilog HDL does not have library for any the mathematical function? So, I am hopefully received your detailed steps ?

Thanks so much!



Adrian Rosebrock November 14, 2016 at 12:09 pm #

REPLY ↩

Hi Cam — I can't speak for the FGPA side of things (since I don't have experience in that area), but I provided *very detailed* steps + source code using the HOG + Linear SVM method [inside the](#)

Free 21-day crash course
on computer vision &
image search engines

×



Yash Karnawat November 14, 2016 at 7:46 am #

Hi Adrian ,

I am a huge fan of your blog and really appreciate w for sign language detection . I am facing some probl image. The hand can be in a square box of size in th

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

apply both the image pyramid technique and hard-negative mining to train my binary classifier and how do I go about the retraining process?



Adrian Rosebrock November 14, 2016 at 12:00 pm

The actual size of the hand doesn't matter as long as the aspect ratio (ratio of width to height) is the same. That is the entire point of applying a sliding window + image pyramid — to detect objects at various scales and locations of an image. I detail a thorough solution to training a custom HOG + Linear SVM object detector [inside the PyImageSearch Gurus course](#). I would suggest starting there.

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Saloni Mittal November 22, 2016 at 10:41 am #

REPLY ↩

Hi Adrian,

How can we retrain an existing svm without starting from scratch the retraining process?

While performing hard negative training I got a Memory Error in the end, apparently the number of final hog features exceeded the numpy array size! (My dataset contains 10K positives and 60K negatives, but I performed hard neg mining on 16K negatives. The number of hog features extracted from every image is 1568).

Thanks in advance! 😊

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock November 22, 2016 at 12:21 pm #

REPLY ↩

There are machine learning algorithms that can be trained in batches and sequentially updated — SVMs are not one of them. In this case, you'll need to reduce the size of your training dataset. I would suggest reducing the size of the true negatives and keeping only the most confident hard-negatives.

Trackbacks/Pingbacks

[Non-Maximum Suppression for Object Detection in Python](#)
 [...] If you remember, last week we discussed [Histogram of Oriented Gradients](#)
[Capturing mouse click events with Python and OpenCV](#)
 [...] In this example we'll click and drag a rectangular Region of Interest (ROI) on the image.
 This technique is especially helpful if you are labeling data for training a classifier.
[Image Pyramids with Python and OpenCV - PyImageSearch](#)
 [...] see, a few months ago I wrote a blog post on utilizing the HOG
 descriptor and a Linear SVM to detect objects in images
[Sliding Windows for Object Detection with Python and OpenCV](#)

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

[...] fact, both sliding windows and image pyramids are classification [...]

[Pedestrian Detection OpenCV - PyImageSearch](#) - November 11, 2015

[...] OpenCV ships with a pre-trained HOG + Linear SVM detector in both images and video streams. If you're not familiar with the HOG + Linear SVM method, I suggest you read this blog post:

[HOG detectMultiScale parameters explained - PyImageSearch](#)

[...] accomplish this, we leveraged the built-in HOG + Linear SVM detector that OpenCV ships with, allowing us to detect people in [...]

[Detecting machine-readable zones in passport images - PyImageSearch](#) - November 30, 2015

[...] to detect the MRZ region of a passport that you need a bit of machine learning, perhaps using the Linear SVM + HOG framework to construct an "MRZ detector" — but that would be [...]

[Intersection over Union \(IoU\) for object detection - PyImageSearch](#) - November 7, 2016

[...] is interested in building a custom object detector using the HOG + Linear SVM framework for his final year project. He understands the steps required to build the object detector well [...]

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Leave a Reply

Free 21-day crash course on
computer vision & image search
engines

SUBMIT COMMENT

Resource Guide (it's totally free).



Click the button below to get my **free 11-page** **Image Search Engine Resource Guide**. Uncover **exclusive techniques** that I don't teach elsewhere to build your own image search engines of your own.

Free 21-day crash course
on computer vision &
image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to

Download

Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!


KICKSTARTER: Deep Learning for Computer Vision

KICKSTARTER

DEEP LEARNING
FOR COMPUTER VISION

WITH PYTHON

Dr. Adrian Rosebrock




Free 21-day crash course on
computer vision & image search
engines



You're interested in deep learning and computer vision, *but you don't know how to get started*. Let me help. [My new book will teach you all you need to know about deep learning.](#)

CLICK HERE TO CHECK OUT THE KICKSTARTER

You can detect faces in images & video.




Are you interested in **detecting faces in images & video**? Let me help! I guarantee that my new book will turn you into [here to give it a shot yourself](#).

Free 21-day crash course
on computer vision &
image search engines



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

[CLICK HERE TO MASTER FACE DETECTION](#)**PyImageSearch Gurus: NOW ENROLLING!****The PyImageSearch Gurus course is *now enrolling!* Inside [Free 21-day crash course on](#)**

- Automatic License Plate Recognition (ANPR)
- Deep Learning
- Face Recognition
- *and much more!*


[computer vision & image search engines](#)**Click the button below to learn more about the course, take a tour, and get 10 (FREE) sample lessons.**[TAKE A TOUR & GET 10 \(FREE\) LESSONS](#)**Hello! I'm Adrian Rosebrock.**

I'm an entrepreneur and Ph.D who has launched two successful image search engines, [ID My Pill](#) and [Chic Engine](#). I'm here to share my tips, tricks, and hacks I've learned along the way.

Learn computer vision in a single weekend.**Free 21-day crash course
on computer vision &
image search engines**

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good?

Enter your email below to start your journey to



Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Want to learn computer vision & OpenCV? I can teach you in a **single weekend**. I know. It sounds crazy, but it's no joke. My new book is your **guaranteed, quick-start guide** to becoming an OpenCV Ninja. So why not give it a try? [Click here to become a computer vision ninja.](#)

CLICK HERE TO BECOME AN OPENCV NINJA

Subscribe via RSS



Never miss a post! Subscribe to the PyImageSearch RSS Feed and keep up to date with my latest search engine tutorials, tips, and tricks

Free 21-day crash course on
computer vision & image search
engines



POPULAR

Install OpenCV and Python on your Raspberry Pi 2 and B+

FEBRUARY 23, 2015

Home surveillance and motion detection with the Raspberry Pi, Python, OpenCV, and Dropbox

JUNE 1, 2015

How to install OpenCV 3 on Raspbian Jessie

OCTOBER 26, 2015

Install OpenCV 3.0 and Python 2.7+ on Ubuntu

JUNE 22, 2015

Basic motion detection and tracking with Python and OpenCV

MAY 25, 2015

Install OpenCV 3.0 and Python 2.7+ on OSX

JUNE 15, 2015

Accessing the Raspberry Pi Camera with OpenCV and Python

MARCH 30, 2015

Free 21-day crash course
on computer vision &
image search engines



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to

Search

Search...

Enter your email below to start your journey to becoming a computer vision master.

Email Address

LET'S DO IT!

Find me on [Twitter](#), [Facebook](#), [Google+](#), and [LinkedIn](#).

© 2017 PyImageSearch. All Rights Reserved.

Free 21-day crash course on
computer vision & image search
engines



Free 21-day crash course on computer vision & image search engines



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to